

Numerical mathematics on FPGAs using CλaSH

Comparison to conventional solution methods

Martijn Bakker

Committee:

Dr.ir. J. Kuper

Dr. R.M.J. van Damme

Dr.ir. J. Broenink

Computer Architecture for Embedded Systems Electrical Engineering,
Mathematics and Computer Science (EEMCS)

University of Twente

This thesis is submitted for the degree of

Bachelor of Science

Advanced Technology

July 2nd, 2015

Abstract

This is where you write your abstract ...

Acknowledgements

- Jan Kuper
- Christiaan Baaij - making CλaSH and answering questions
- Ruud van Damme
- Jan Broenink

Acronyms

ASIC	Application-specific integrated circuit
FPGA	Field-Programmable Gate Array
CPU	Central Processing Unit
GPU	Graphics Processing Unit
GPGPU	General Purpose GPU
DSP	Digital Signal Processing
LED	Light Emitting Diode
VHDL	VHSIC HDL
VHSIC	Very High Speed Integrated Circuit
HDL	Hardware Description Language
SoC	System-on-Chip
ODE	Ordinary Differential Equation

Table of contents

1	Introduction	1
1.1	Project goals	1
1.2	FPGAs	1
1.2.1	What is an FPGA?	1
1.2.2	How does it work?	2
1.2.3	System-on-a-chip	3
1.3	Numerical solvers for ODEs	3
1.4	Functional programming	5
1.4.1	What is functional programming?	5
1.4.2	Using FP for numerical mathematics	6
1.4.3	Example: Numerical solutions of ODEs in Haskell	7
1.5	CλaSH	9
1.5.1	Designing hardware in a functional way	9
1.5.2	Mealy machines	9
2	Results	11
2.1	Euler's method	11
2.1.1	Performance	11
2.1.2	Accuracy	11
2.2	Second order Runge-Kutta	11
2.2.1	Performance	11
2.2.2	Accuracy	11
2.3	4th order Runge-Kutta	11
2.3.1	Performance	11
2.3.2	Accuracy	11
2.4	Comparison with CPU implementations	11
2.5	Comparison with GPU implementations	11
3	Conclusion	13
4	Discussion	15
4.1	sdfasdf	15
	References	17

Appendix A	Haskell source code for numerical solutions of ODEs	19
Appendix B	Installing the CUED class file	25
Appendix C	Installing the CUED class file	27
Appendix D	Installing the CUED class file	29

Chapter 1

Introduction

This introduction contains the project goals and some very short introductions on the FPGAs and how they are used, numerical methods for approximating ODEs, functional programming and C λ aSH. The short introductions are included for the sake of completeness and are assumed knowledge for the rest of the thesis.

1.1 Project goals

From the start, this project has had two main goals: firstly, obtaining information on the feasibility and the advantages and disadvantages of performing numerical mathematics (numerical approximations to ODEs) directly on (programmable) hardware, the FPGA. Secondly: figuring out whether higher-order functions are of much use for numerical mathematics. As per usual, having main goals spawns off several minor goals which support the main parts. Both supporting goals are about simplifying the process of configuring FPGAs: an easy way of setting up projects with complicated IO requirements and furthermore, developing a tool-chain integration which turns the long process of compiling and deploying your FPGA project into the execution of a single command.

Alongside these concrete goals the underlying theme is to do as much work as possible in C λ aSH, a library and compiler based on the functional programming language Haskell, developed by Christiaan Baaij at the CAES group at the University of Twente. Further elaboration on C λ aSH can be found in section 1.5.

1.2 FPGAs

1.2.1 What is an FPGA?

An FPGA (Field Programmable Gate Array) is a chip in which you can specify the hardware yourself. In contrast to regular programming in which you generate a long list of instructions which are executed sequentially on a fixed chip configuration, the FPGA allows you to specify exactly which wire (signal) leads where and what operation should be applied to that signal. This approach to programming can have several advantages. The first one arises

from the large opportunities for parallelism. Every part of the FPGA can be executing a meaningful computation simultaneously, whereas processors are bound by the amount of physical cores they have in the amount of truly concurrent instruction executions possible. Secondly, a conventional processor only has a fixed instruction set. Using an FPGA you can define your own instructions (subcircuits), again providing a possible improvement in computational speed. According to [6], FPGAs were already capable of outperforming CPUs on very parallelizable numerical tasks on single and double precision floating point numbers in as early as 2004. Furthermore, as you are implementing your signal processing directly in hardware, there will be a fixed bound on the possible latency. This makes FPGAs ideal for purposes in high-throughput, low-latency signal processing, eg. real-time audio, video or data stream processing. Lastly, the reconfigurability of FPGAs whilst remaining close to the actual hardware allows for cost reductions in the verification of ASIC (Application-Specific Integrated Circuits) designs. It's cheaper to reprogram your FPGA than to have a new version of an ASIC manufactured.

However, the FPGA is a trade-off between implementing designs directly in hardware and being able to run multiple designs (after a reconfiguration). As a consequence of this, it still loses to ASICs with several orders of magnitude on performance [2] and CPUs still dominate in terms of versatility and on-the-fly reconfigurability.

1.2.2 How does it work?

An FPGA is built up from several distinct element types, depicted in figure 1.1

1. *Logic elements* Responsible for the actual signal processing. An FPGA may contain different types of logic elements, eg. memory, DSP and logic blocks. These blocks implement some signal processing capability which can be configured up to certain limits.
2. *Programmable interconnects* In order to be able to represent complex designs, the logic elements need to be connected in a certain way. This is what the programmable interconnects are for. Essentially, those are wires which can be turned on and off by the user as part of a design specification.
3. *IO blocks* Finally, the functionality implemented using the logic elements and programmable interconnects should be exposed to external signals in order to be useful. IO blocks can be used to control hardware pins, controlling a LED or reading a switch but also for more intricate IO facilities, eg. external memory controllers.

It might seem that programming an FPGA involves manually specifying the interconnects and exact configurations of the logic elements. However, specialized languages have been developed just for the purpose of describing the FPGA functionality at a higher level of abstraction and leave the specific routing and assignment of logic elements to the compiler. The two main advantages of these languages are that you do not have to worry about low-level problems like how the interconnects will be routed and secondly, your written specification will be portable across multiple FPGA vendors as long as the vendor supplies you with the proper compiler from your specification to a file which can be used to program the FPGA.

More information on FPGA functionality can be found in [7].




Fig. 1.1 FPGA fabric.

1.2.3 System-on-a-chip

FPGAs in itself can be useful, but especially for design with IO requirements that are more complex than just reading out hardware switches and controlling LEDs, more control is needed. This requirement has led to the rise of SoCs (System-on-Chip). These devices integrate an FPGA with additional hardware on a single chip. This extra hardware usually contains a CPU, which can be used to simplify the process of loading and extracting data from the FPGA. The SoC used for this thesis is the Terasic SoCKit, a development kit containing an Altera Cyclone V FPGA and a dual-core ARM A9 CPU on a single chip (Altera 5CSXFC6D6F31C6N), alongside a wide variety of IO possibilities. Further information on the SoC used is available at [3].

1.3 Numerical solvers for ODEs

The field of numerical solvers is a vast and active area of research. Furthermore, there is a lot of theory on which solver to pick for specific problems, related to stability, computational efficiency and other factors. However, the goal of this thesis is to get an impression of the feasibility of using numerical solvers for ODEs on FPGAs. Therefore, the selection of solvers will be restricted to some very basic schemes.

The solvers used have some common properties:

- *Fixed-step* The solvers compute a value for the ODE at a fixed step size. This means that in contrary to the continuous 'mathematical' solution of an ODE, the output of the solver will be an approximate to the actual value of the solution at the discrete set of values of the independent variable, in which the difference between the values of the independent variables is defined by the fixed step size.
- *Single-step* The approximate value of the ODE x_k at t_k only depends on t_{k-1} , x_{k-1} and the ODE that is being approximated.

Euler

The easiest numerical solver, without doubt, is Euler's method (equation 1.1). For every discrete point in time, the value of the next point is approximately equal to the derivative at that point multiplied by the time step added to the current value. The simplicity of the scheme has a cost: it is not very accurate and the errors accumulate quickly. From [5], the maximum error of Euler's method can be shown to be linear in the time step and exponential in the interval length (eq 1.2), in which M and L are constants depending on the equation to be solved, $b - a$ is the solution interval length and h is the time step.

$$\begin{aligned} s_{k,1} &= f(t_k, x_k) \\ x_{k+1} &= x_k + h s_{k,1} \end{aligned} \tag{1.1}$$

$$\text{maximum error}_{\text{euler}} \leq \frac{Mh}{L} (e^{L(b-a)} - 1) \tag{1.2}$$

Runge-Kutta methods

The Runge-Kutta methods are a family of solvers, of which the 4th order version is the most well-known (RK4). The solver used here will be a second-order Runge-Kutta method, also known as the *improved Euler's method* [5]. This second order method requires the computation of two slopes (equation 1.3), in contrast to the single one required for Euler's method.

$$\begin{aligned} s_{k,1} &= f(t_k, x_k) \\ s_{k,2} &= f(t_k + h, x_k + h s_{k,1}) \\ x_{k+1} &= x_k + h \frac{s_{k,1} + s_{k,2}}{2} \end{aligned} \tag{1.3}$$

$$\text{maximum error}_{\text{RK2}} \leq \frac{Mh^2}{L} (e^{L(b-a)} - 1) \tag{1.4}$$

Note that the maximum error of RK2 is proportional to the square of the time step (equation 1.4). As the time step decreases by a factor of 2, the maximum error will decrease by a factor of 4, given that the equation and the range stay the same. However, the maximum error still depends exponentially on the interval length.

1.4 Functional programming

1.4.1 What is functional programming?

As the name suggests, the functional programming paradigm uses functions. These functions are used to build up the program and create structure. In contrast to the imperative programming paradigm, there is no assignment - there are only expressions. It is by evaluation of these expressions that you execute your program. These expressions consist of variables, constants and operations. However, the name variable may be ill-chosen, as assignment does not exist and therefore it is impossible for a variable to vary. Once you have bound a value to a certain variable, this value may not change and the value should be the same at every point where this variable is referenced (a concept called referential transparency).

The exclusion of assignment from functional languages has several consequences. It becomes impossible to program using loops. The alternative is the use of recursive (listing ??) and higher-order functions (functions that have other functions as parameters or output). Especially higher-order functions have the side effect that they add clarity about the way the program functions (listing 1.2). For instance, a **map** applies the same function to every element in a list, resulting in a new list. A **fold** would start at an initial value and element-by-element combine the list into a single value using a specified function, eg. addition. Both these operations would be implemented using a for-loop in an imperative language, but their goal is completely different. The availability of higher-order functions allows for more clarity in code by being able to specify exactly what kind of operation you want to perform. Another effect of the lack of assignment is the lack of side effects in functional programming. As there is no way to modify a variable, there is also no way of accidentally modifying a variable such that you enter an invalid state and other parts of the program stop functioning correctly.

Listing 1.1 Recursive functions

```
1 fac :: Num a => a -> a
2 fac 0 = 1
3 fac n = n * fac(n-1)
```

Listing 1.2 Higher order functions

```
1 timesTwo :: Num a => [a] -> [a]
2 timesTwo xs = map (*2) xs
3
4 sum :: Num a => [a] -> a
5 sum xs = foldl (+) 0 xs
```

The concept of higher-order functions also serves as an introduction to another very important concept in Haskell: the type system. Everything has a fixed type and often times, when only looking at the type definition you can already guess what the function is going to do. Consider the type signature of **map** in listing 1.3. It requires a $(a \rightarrow b)$, a function to turn something of type a into something of type b . Furthermore, it needs a list of a , $[a]$ (indicated by the square brackets) and it returns something of type 'list of b ' ($[b]$). The type signature of the **foldl** is slightly harder to understand, but it requires a function which needs an a and a b to produce a new b . Furthermore, an initial value (a) and a list of b to operate on ($[b]$) in order

to return the final result, which has again type `a`. Function type signatures can be daunting to understand at first, but not all of them are as complicated as the `foldl`. For instance, a function which can be used to represent a differential equation, which needs a state of the system (the `ODEState`) in order to compute the derivative at that point (the `D_ODEState`). As a last example, the type signature of `Solver`: it requires some integration scheme, settings for the time (initial time and a time step), it requires an equation and an initial state of the system. All of this combined results in a list of states: the numerical approximation to the solution of the ODE. Note that in this example, the integration scheme `Scheme` itself is a function, for which understanding the type signature should pose no problem by now.

Listing 1.3 Type signatures

```

1 map :: (a->b) -> [a] -> [b]
2 foldl :: (a->b->a) -> a -> [b] -> a
3
4 type Equation = ODEState -> D_ODEState
5 type Scheme   = TimeSettings -> Equation -> ODEState -> ODEState
6 type Solver   = Scheme -> TimeSettings -> Equation -> ODEState -> [ODEState]
```

1.4.2 Using FP for numerical mathematics

Functional languages have several properties which make them suitable for the purpose of solving problems in numerical mathematics. First and foremost, Haskell, being based on λ -calculus is very close to mathematics. The useful mathematical properties here are *referential transparency*, *easy partial function application* and being a *declarative language*. Referential transparency implies that a variable only has a constant value which is the same everywhere in the program. This prevents that changing a variable might have influence on another computation as a side effect and it corresponds to mathematical notation. For instance, in an imperative programming like C you could write `i = i + 1`, which is a mathematical impossibility and therefore not allowed in Haskell. Partial function application is another very useful concept. Often in numerical mathematics, you want to create or process a function. You need a function that has another function as return value. For instance, take a function which requires two arguments. After only applying a single argument, the object returned still needs the second argument in order to compute the final value. This is exactly according to the definition of a function: An object that still needs arguments before being able to return its final value. Being a *declarative language* means that you write code that specifies what you want to accomplish, not how to get there. This concept is again borrowed from mathematics. You put in a set of function definitions and Haskell will figure out how to actually compute the value you request according to those definitions. This property of declarativity also has the result that Haskell is a terse language whilst remaining easy to understand. Secondly, Haskell has a very strong type system. The type system has three main advantages. It becomes very easy to swap out and replace functions as long as you make sure that the types are the same. The Haskell compiler will start to assert errors immediately whenever you feed it something which does not make sense or could be ambiguous which is very useful when writing programs. By having a look at the types of a Haskell program it

becomes very straightforward to see what the program does and how it works, which is very useful when attempting to understand your own or someone else's code. Lastly, a property which is often very important for numerical mathematics: Haskell is fast. According to the *Computer Language Benchmarks Game* [1], Haskell is almost on par with Java and Fortran but significantly faster than Python and Matlab (not shown), two languages which are often used for numerical mathematics nowadays. There is still a performance gap of around a factor 3 between Haskell and C (the reference), hence if speed is of the absolute highest concern C is still a valid option.

1.4.3 Example: Numerical solutions of ODEs in Haskell

As mentioned before, the types in Haskell reveal lots of information about the structure and functionality of the program. The three main types constituting the numerical solver for ordinary differential equations are listed above.

Listing 1.4 Main types for the ODE solver

```
1  type Equation = ODEState -> D_ODEState
2  type Scheme   = TimeSettings -> Equation -> ODEState -> ODEState
3  type Solver    = Scheme -> TimeSettings -> Equation -> ODEState -> [ODEState]
```

Equation

In essence, a differential equation is a mapping (function) from a certain state of the system to the change of this system. This is also what the type signature of Equation signifies, a mapping from an ODEState to a D_ODEState. This generic set up allows the specification of any ODE for solving. The implementation in pure Haskell of a simple ODE is given in listing 1.5, which corresponds to the equation $x' = -x$. However, this representation is not very elegant and a lot of the code is performing unboxing of the types. Using property that this equation is linear, it is possible to use an utility method which takes as input a matrix and returns the Haskell differential equation function belonging to that matrix. The same can be done for heterogeneous linear systems using a different utility function, which does not only takes a matrix as input but also a list of functions representing the heterogeneous part of the equation. The example code for this can be seen in appendix A.

Listing 1.5 Example equation for exponential decay

```
1  eq_exponential :: Equation
2  eq_exponential state = [-x !! 0]
3  where
4    x = xs state
```

SolveMethod

The SolveMethod performs the actual computations on what the next value of the solution should be: the integration scheme. In order to obtain this next state, the scheme needs three input values: It needs information on the timing constraints of the solution, in this case it

needs the time step. Furthermore, it needs the equation itself and it requires the state of the system at t_n in order to be able to determine the state of the system at $t_{n+1} = t_n + \Delta t$.

The most straightforward integration scheme is called forward Euler, given in equation 1.1. Listing ?? depicts the translation of the mathematical expression 1.1 to Haskell. Even though some list operations have been inserted (**zipWith** and **map**), the structure is still recognizable. It computes the change in state, multiplies this with the time step obtained in line 6 and adds the initial state in line 4. Lastly, the integration scheme returns the new state of the equation, consisting of a list of x-values and a corresponding time value. Implementations of different solvers (eg. 4th order Runge-Kutta) can be found in appendix A.

Listing 1.6 Example code for the Euler integration scheme

```

1  euler :: Scheme
2  euler time equation initState = ODEState newX newT
3      where
4      newX      = zipWith (+) (xs initState) dX
5      dX        = map (timestep *) (equation initState)
6      newT      = (t initState) + timestep
7      timestep  = dt time

```

Solver

The Solver function in listing 1.7 acts as the main interface to the program. You specify a SolveMethod, the TimeSettings (containing the time step and the time at which to stop solving), the equation itself and an initial condition. The Solver will then return a list of states of the system. As is very common in functional programming, the Solver has been defined recursively. Line 4 is where the magic happens: the solution list is defined to be the initial condition, followed by the solution list with the new state (computed by the integration scheme on line 6) as initial condition. Additionally, there is a comparison in line 7 which ends the recursion whenever the time of the solution exceeds the maximum time value, set in the TimeSettings.

The solutions of a wide range of equations, both linear and non-linear, both homogeneous and heterogeneous and using the input matrix utility functions have been plot with suitable initial conditions to show their behavior in figure 1.2.

Listing 1.7 The main controlling function

```

1  solve :: Solver
2  solve solvemethod time equation initState
3      | end = []
4      | otherwise = initState : solve solvemethod time equation newState
5      where
6      newState = solvemethod time equation initState
7      end = (t initState) > (tMax time)

```

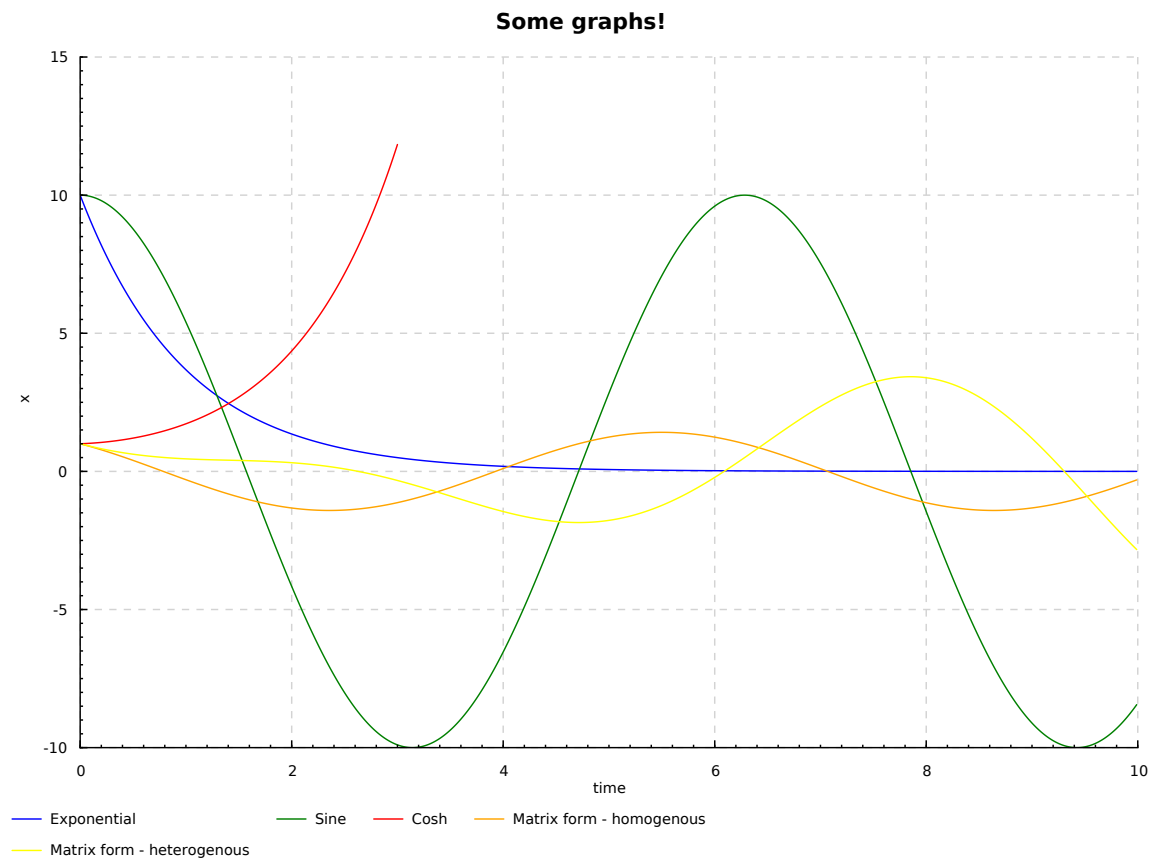


Fig. 1.2 Graphs

Exponential	$x(t)' = -x(t)$	(1.5)
Simple harmonic	$x(t)'' = -x(t)$	(1.6)
Cosine hyperbolic	$x(t)' = \frac{\sqrt{x(t)^2 - a^2}}{a}$	(1.7)
Simple harmonic	$\vec{x}(t)' = \begin{bmatrix} 0 & 1 \\ -1 & 0 \end{bmatrix} \vec{x}(t)$	(1.8)
Simple forced harmonic	$\vec{x}(t)' = \begin{bmatrix} 0 & 1 \\ -1 & 0 \end{bmatrix} \vec{x}(t) + \begin{bmatrix} \sin(t) \\ e^{-t} \end{bmatrix}$	(1.9)

1.5 CλaSH

1.5.1 Designing hardware in a functional way

1.5.2 Mealy machines

Chapter 2

Results

2.1 Euler's method

2.1.1 Performance

2.1.2 Accuracy

2.2 Second order Runge-Kutta

2.2.1 Performance

2.2.2 Accuracy

2.3 4th order Runge-Kutta

2.3.1 Performance

2.3.2 Accuracy

2.4 Comparison with CPU implementations

2.5 Comparison with GPU implementations

Chapter 3

Conclusion

Chapter 4

Discussion

4.1 sdfasdf

References

- [1] Computer language benchmarks game. <http://benchmarksgame.alioth.debian.org/u32q/which-programs-are-fastest.html>, 2015. [Accessed: June 2015].
- [2] Mining hardware comparison. https://en.bitcoin.it/wiki/Mining_hardware_comparison, 2015. [Accessed: June 2015].
- [3] Sockit - the development kit for new soc device. <http://www.terasic.com.tw/cgi-bin/page/archive.pl?Language=English&CategoryNo=&No=816>, 2015. [Accessed: June 2015].
- [4] J Kuper. *A Short Introduction to Functional Programming*. 2015.
- [5] J Polking, A Boggess, and D Arnold. *Differential Equations with Boundary Value Problems*. Prentice Hall, Upper Saddle River, New Jersey, 2006.
- [6] Keith Underwood. Fpgas vs. cpus: Trends in peak floating-point performance. *Association for Computing Machinery*, 2004. URL <http://home.engineering.iastate.edu/~zambreno/classes/cpre583/2006/documents/Und04A.pdf>.
- [7] Wayne Wolf. *FPGA-Based System Design*. Prentice Hall, Upper Saddle River, New Jersey, 2004.

Appendix A

Haskell source code for numerical solutions of ODEs

Listing A.1 Solver.hs

```
1
2 module Solver where
3
4 import Prelude
5 import SolverTypes
6
7 import SolverEquations
8 import SolverSolvers
9 import SolverPlotter
10 import SolverPresets
11
12 ----- CALLERS
13 -- general form, stop after a certain time
14 solve :: Solver
15 solve solvemethod time equation initState
16   | end = []
17   | otherwise = initState : solve solvemethod time equation newState
18 where
19   newState = solvemethod time equation initState
20   end = (t initState) > (tMax time)
21
22 sol_start = solve rk4 initTimeSettings
23
24 solution_expo = sol_start eq_exponential initODEState
25 solution_sine = sol_start eq_sine initODEState
26 solution_cosh = solve rk4 initTimeSettings2 eq_cosh initODEState2
27 solution_homo = sol_start (eq_linear_homo_const sinematrix) initODEState2
28 solution_hetr = sol_start (eq_linear_hetr_const sinematrix funcvec) initODEState2
29
30
31 testPlot = plotSolutions [s1,s2,s3,s4,s5] "Some graphs!"
32 where
33   s1 = (solution_expo, "Exponential")
```

```

34     s2 = ( solution_sine , "Sine")
35     s3 = ( solution_cosh , "Cosh")
36     s4 = ( solution_homo, "Matrix form – homogenous")
37     s5 = ( solution_hetr , "Matrix form – heterogenous")

```

Listing A.2 SolverTypes.hs

```

1  module SolverTypes where
2
3  import Prelude
4
5  type NumRepr = Float
6  type D_ODEState = [NumRepr]
7
8  data ODEState = ODEState { xs :: [NumRepr]
9    , t :: NumRepr
10 } deriving (Show)
11
12 data TimeSettings = TimeSettings { dt :: NumRepr
13   , tMax :: NumRepr
14 } deriving (Show)
15
16 type SubFunction = (NumRepr -> NumRepr)
17
18 type Equation = ODEState -> D_ODEState
19 type Scheme   = TimeSettings -> Equation -> ODEState -> ODEState
20 type Solver   = Scheme -> TimeSettings -> Equation -> ODEState -> [ODEState]

```

Listing A.3 SolverEquations.hs

```

1  module SolverEquations where
2
3  import Prelude
4  import SolverTypes
5  import SolverPresets
6
7  -- Exponential :  $y = A * \exp(-t)$ 
8  --  $y' = -y$ 
9
10 --  $x0' = -x0$ 
11 eq_exponential :: Equation
12 eq_exponential state = [-x !! 0]
13   where
14     x = xs state
15
16
17
18 -- Sine :  $y = A * \sin(\omega * t)$ 
19 --  $y'' = -y$ 
20
21 --  $x0' = x1$ 
22 --  $x1' = -x0$ 
23 eq_sine :: Equation

```

```

24 eq_sine state    = [x0,x1]
25   where
26     x = xs state
27     x0 = x !! 1
28     x1 = - (x !! 0)
29
30
31
32 -- Hyperbolic cosine :  $y = a \cosh((x - b)/a)$ 
33 --  $y' = \sqrt{y^2 - a^2}/a$ 
34
35 --  $x0' = \sqrt{x0^2 - a^2}/a$ 
36 eq_cosh :: Equation
37 eq_cosh state = [x0]
38   where
39     x = xs state
40     x0 = sqrt((x !! 0)^2 - a^2)/a
41     a = 0.99
42
43 -- Arbitrary homogenous system
44 --  $y' = Ay$ 
45 eq_linear_homo_const :: [[NumRepr]] -> ODEState -> D_ODEState
46 eq_linear_homo_const matrix state = map (rowmult y) matrix
47   where
48     y = xs state
49
50 rowmult :: [NumRepr] -> [NumRepr] -> NumRepr
51 rowmult vec1 vec2 = sum $ zipWith (*) vec1 vec2
52
53 -- Arbitrary heterogenous system
54 --  $y' = Ay + F$ 
55 eq_linear_hetr_const :: [[NumRepr]] -> [SubFunction] -> ODEState -> D_ODEState
56 eq_linear_hetr_const matrix vector state = zipWith (+) (map (rowmult y) matrix) (map ($time) vector)
57   where
58     y = xs state
59     time = t state

```

Listing A.4 SolverPresets.hs

```

1 module SolverPresets where
2
3 import Prelude
4 import SolverTypes
5
6 unity :: [[NumRepr]]
7 unity = [[1,0,0],[0,1,0],[0,0,1]]
8
9 sinematrix :: [[NumRepr]]
10 sinematrix = [[0,1],[-1,0]];
11
12 vec :: [NumRepr]
13 vec = [4,3,2]
14

```

```

15 funcvec :: [SubFunction]
16 funcvec = [(\ t -> sin t ),(\ t -> exp (-t))]
17
18 initODEState = ODEState [10, 0.0] 0.0
19 initODEState2 = ODEState [1, -1] 0.0
20
21 initTimeSettings = TimeSettings 0.01 10
22 initTimeSettings2 = TimeSettings 0.01 3

```

Listing A.5 SolverHelper.hs

```

1 module SolverHelper where
2
3 import Prelude
4 import SolverTypes
5
6 sumLists :: [[NumRepr]] -> [NumRepr] -> [NumRepr]
7 sumLists [] factors = []
8 sumLists (xs:[]) factors = map ((head factors)*) xs
9 sumLists (xs:xss) factors = zipWith (+) (map (head factors*) xs) (sumLists xss (tail factors))

```

Listing A.6 SolverPlotter.hs

```

1 module SolverPlotter where
2
3 import Prelude
4 import SolverTypes
5 import GHC.Float
6
7 import Graphics.Rendering.Chart.Easy
8 import Graphics.Rendering.Chart.Backend.Cairo
9
10 outProps = fo_format .~ PDF $ def
11
12 plotSolutions :: [(ODEState, String)] -> String -> String -> IO()
13 plotSolutions solutions title filename = toFile outProps filename $ do
14   layout_title .= title
15   layout_x_axis . laxis_title .= "time"
16   layout_y_axis . laxis_title .= "x"
17   plotSolutions_help solutions
18
19
20 plotSolutions_help [] = error "empty list "
21 plotSolutions_help [sol] = plotSolution sol
22 plotSolutions_help (sol:sols) = do
23   plotSolution sol
24   plotSolutions_help sols
25
26
27 plotSolution (solution, curveTitle) = plot $ line curveTitle [states]
28 where
29   states = reformData solution
30

```

```
31
32 reformData :: [ODEState] -> [(Double, Double)]
33 reformData states = map reformState states
34
35 reformState :: ODEState -> (Double, Double)
36 reformState state = (float2Double tVal, float2Double (x !! 0))
37   where
38     x = xs state
39     tVal = t state
```


Appendix B

Installing the CUED class file

\LaTeX .cls files can be accessed system-wide when they are placed in the $\langle\text{texmf}\rangle/\text{tex}/\text{latex}$ directory, where $\langle\text{texmf}\rangle$ is the root directory of the user's \TeX installation. On systems that have a local texmf tree ($\langle\text{texmflocal}\rangle$), which may be named “ texmf-local ” or “ localtexmf ”, it may be advisable to install packages in $\langle\text{texmflocal}\rangle$, rather than $\langle\text{texmf}\rangle$ as the contents of the former, unlike that of the latter, are preserved after the \LaTeX system is reinstalled and/or upgraded.

It is recommended that the user create a subdirectory $\langle\text{texmf}\rangle/\text{tex}/\text{latex}/\text{CUED}$ for all CUED related \LaTeX class and package files. On some \LaTeX systems, the directory look-up tables will need to be refreshed after making additions or deletions to the system files. For \TeX Live systems this is accomplished via executing “ texhash ” as root. $\text{MIK}\text{\TeX}$ users can run “ initexmf -u ” to accomplish the same thing.

Users not willing or able to install the files system-wide can install them in their personal directories, but will then have to provide the path (full or relative) in addition to the filename when referring to them in \LaTeX .

Appendix C

Installing the CUED class file

\LaTeX .cls files can be accessed system-wide when they are placed in the $\langle\text{texmf}\rangle/\text{tex}/\text{latex}$ directory, where $\langle\text{texmf}\rangle$ is the root directory of the user's \TeX installation. On systems that have a local texmf tree ($\langle\text{texmflocal}\rangle$), which may be named “ texmf-local ” or “ localtexmf ”, it may be advisable to install packages in $\langle\text{texmflocal}\rangle$, rather than $\langle\text{texmf}\rangle$ as the contents of the former, unlike that of the latter, are preserved after the \LaTeX system is reinstalled and/or upgraded.

It is recommended that the user create a subdirectory $\langle\text{texmf}\rangle/\text{tex}/\text{latex}/\text{CUED}$ for all CUED related \LaTeX class and package files. On some \LaTeX systems, the directory look-up tables will need to be refreshed after making additions or deletions to the system files. For \TeX Live systems this is accomplished via executing “ texhash ” as root. \TeX users can run “ initexmf -u ” to accomplish the same thing.

Users not willing or able to install the files system-wide can install them in their personal directories, but will then have to provide the path (full or relative) in addition to the filename when referring to them in \LaTeX .

Appendix D

Installing the CUED class file

\LaTeX .cls files can be accessed system-wide when they are placed in the $\langle\text{texmf}\rangle/\text{tex}/\text{latex}$ directory, where $\langle\text{texmf}\rangle$ is the root directory of the user's \TeX installation. On systems that have a local texmf tree ($\langle\text{texmflocal}\rangle$), which may be named “ texmf-local ” or “ localtexmf ”, it may be advisable to install packages in $\langle\text{texmflocal}\rangle$, rather than $\langle\text{texmf}\rangle$ as the contents of the former, unlike that of the latter, are preserved after the \LaTeX system is reinstalled and/or upgraded.

It is recommended that the user create a subdirectory $\langle\text{texmf}\rangle/\text{tex}/\text{latex}/\text{CUED}$ for all CUED related \LaTeX class and package files. On some \LaTeX systems, the directory look-up tables will need to be refreshed after making additions or deletions to the system files. For \TeX Live systems this is accomplished via executing “ texhash ” as root. \TeX users can run “ initexmf -u ” to accomplish the same thing.

Users not willing or able to install the files system-wide can install them in their personal directories, but will then have to provide the path (full or relative) in addition to the filename when referring to them in \LaTeX .

